

Concurrent Programming

Dr. José Luis Zechinelli Martini

joseluis.zechinelli@udlap.mx

LDS-1101

Based on the course of Professor Carlo A. Furia, Chalmers University of Technology – University of Gothenburg

Agenda

- Concurrent programs
- Races
- Synchronization problems:
 - Mutual exclusion
 - Deadlock
 - Starvation and fairness
- Locks

Abstraction of concurrent programs

- We use an abstract notation for multithreaded applications using Java syntax:

<code>int counter = 0;</code>	
<code>thread t</code>	<code>thread u</code>
<code>int cnt;</code> <code>cnt = counter;</code> <code>counter = cnt + 1;</code>	<code>int cnt;</code> <code>cnt = counter;</code> <code>counter = cnt + 1;</code>

- Each line of code includes exactly one instruction that can be executed atomically

Traces (1)

- A sequence of states gives an execution trace of the concurrent program:

State	t's local	u's local	Shared
1	$pc_t: 1, cnt : \perp$	$pc_u: 1, cnt : \perp$	counter : 0
2	$pc_t: 2, cnt : 0$	$pc_u: 1, cnt : \perp$	counter : 0
3	done	$pc_u: 1, cnt : \perp$	counter : 1
4	done	$pc_u: 2, cnt : 1$	counter : 1
5	done	done	counter : 2

- The program counter pc points to the atomic instruction that will be executed next



Traces (2)

- A sequence of states gives an execution trace of the concurrent program:

State	t's local	u's local	Shared
1	$pc_t: 1, cnt : \perp$	$pc_u: 1, cnt : \perp$	counter : 0
2	$pc_t: 1, cnt : \perp$	$pc_u: 2, cnt : 0$	counter : 0
3	$pc_t: 1, cnt : \perp$	done	counter : 1
4	$pc_t: 2, cnt : 1$	done	counter : 1
5	done	done	counter : 2

- The program counter pc points to the atomic instruction that will be executed next

Traces (3)

- A sequence of states gives an execution trace of the concurrent program:

State	t's local	u's local	Shared
1	$pc_t: 1, cnt : \perp$	$pc_u: 1, cnt : \perp$	counter : 0
2	$pc_t: 2, cnt : 0$	$pc_u: 1, cnt : \perp$	counter : 0
3	$pc_t: 2, cnt : 0$	$pc_u: 2, cnt : 0$	counter : 0
4	$pc_t: 2, cnt : 0$	done	counter : 1
5	done	done	counter : 1

- The program counter pc points to the atomic instruction that will be executed next



Agenda

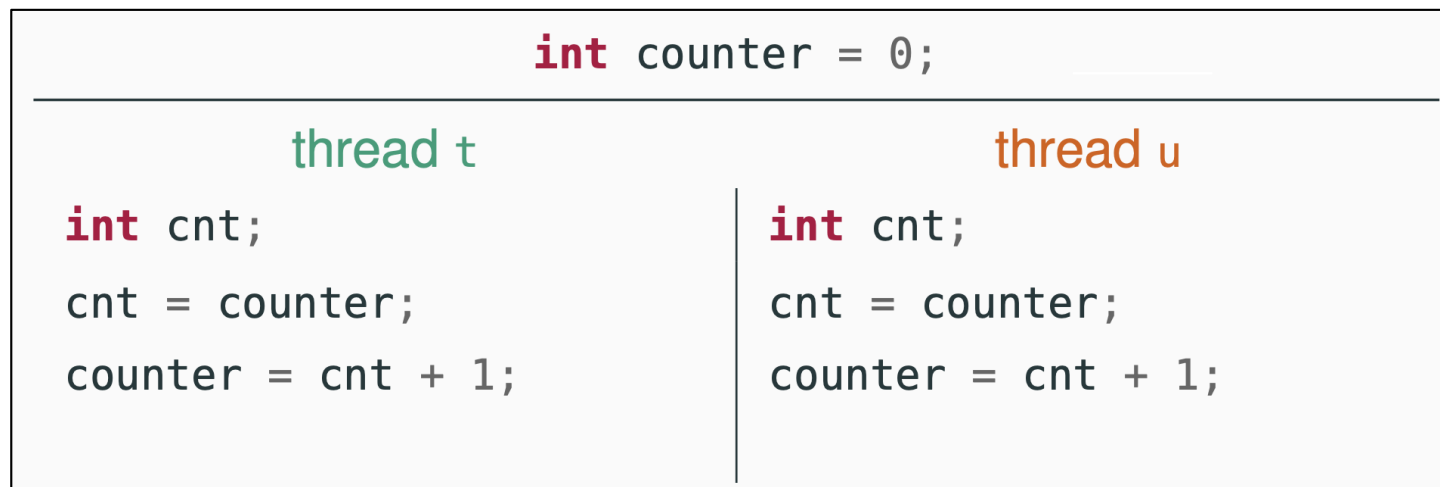
- ✓ Concurrent programs
- Races
- Synchronization problems:
 - Mutual exclusion
 - Deadlock
 - Starvation and fairness
- Locks

Race conditions

- Concurrent programs are **nondeterministic**:
 - Executing multiple times the same concurrent program with the same inputs may lead to different execution traces
 - This is a result of the nondeterministic interleaving of each thread's trace to determine the overall program trace
 - In turn, the interleaving is a result of the scheduler's decisions
- A **race condition** is a situation where the result of a concurrent program depends on the specific execution:
 - The concurrent counter example has a race condition: in some executions the final value of counter is 2, in some executions the final value of counter is 1
 - Race conditions can greatly complicate debugging!

Data races

- Race conditions are typically due lack of synchronization between threads that access shared memory
- A **data race** occurs when two threads access a shared memory location and:
 - At least one access is a write
 - The relative order of the two accesses is not fixed



Data races vs. race conditions

- Not every race condition is a data race:
 - Race conditions can occur even when there is no shared memory access
 - For example in filesystems or network access
- Not every data race is a race condition:
 - The data race may not affect the result
 - For example if two threads write the same value to shared memory

Agenda

- ✓ Concurrent programs
- ✓ Races
- Synchronization problems:
 - Mutual exclusion
 - Deadlock
 - Starvation and fairness
- Locks

Push out the races, bring in the speed

- Concurrent programming introduces:
 - The **potential** for parallel execution (faster, better resource usage)
 - The **risk** of race conditions (incorrect, unpredictable computations)
- The main challenge of concurrent programming is thus introducing parallelism without introducing race conditions:
 - This requires to **restrict** the amount of nondeterminism **by synchronizing** processes/threads that access shared resources

Synchronization

- We will study several synchronization problems that often appear in concurrent programming, together with their solutions:
 - **Correctness** (that is, avoiding race conditions) is more important than performance: an incorrect result that is computed very quickly is no good!
 - However, we also want to retain as much **concurrency as possible**, otherwise we might as well stick with sequential programming



Share memory vs. message passing synchronization (1)

■ Shared memory synchronization:

- Synchronize by writing to and reading from shared memory
- Natural choice in shared memory systems such as threads

■ Message passing synchronization:

- Synchronize by exchanging messages
- Natural choice in distributed memory systems such as processes



Share memory vs. message passing synchronization (2)

- The two synchronization models **overlap**:
 - send a message by writing to and reading from shared memory (example: message board)
 - share information by sending a message (example: order a billboard)
- However, in the first part of the course we will focus on synchronization problems that arise in shared memory concurrency; in the second part we will switch to message passing

The mutual exclusion problem (1)

- The mutual exclusion problem is a fundamental synchronization problem, which arises whenever multiple threads have access to a shared resource:
 - **Critical section**: the part of a program that accesses the shared resource (for example, a shared variable)
 - **Mutual exclusion property**: no more than one thread is in its critical section at any given time
- The mutual exclusion **problem** devise a protocol for accessing a shared resource that satisfies the mutual exclusion property

The mutual exclusion problem (2)

- Simplifications to present solutions in a uniform way:
 - The critical section is an arbitrary **block** of code
 - Threads **continuously** try to enter the critical section
 - Threads spend a **finite amount of time** in the critical section
 - We **ignore what the threads do outside** their critical sections

The mutual exclusion problem (3)

- The **mutual exclusion** problem devise a protocol for accessing a shared resource that satisfies the mutual exclusion property:

T shared;

thread t_j

```
// continuously
while (true) {
    entry protocol
    critical section {
        // access shared data
    }
    exit protocol
} /* ignore behavior
outside critical section */
```

thread t_k

```
// continuously
while (true) {
    entry protocol
    critical section {
        // access shared data
    }
    exit protocol
} /* ignore behavior
outside critical section */
```

The mutual exclusion problem

example: Concurrent counter

- Updating a shared variable consistently is an instance of the mutual exclusion problem:

T shared;

thread t

```
int cnt;  
while (true) {  
    entry protocol  
    critical section {  
        cnt = counter;  
        counter = cnt + 1;  
    }  
    exit protocol  
    return;  
}
```

thread u

```
int cnt;  
while (true) {  
    entry protocol  
    critical section {  
        cnt = counter;  
        counter = cnt + 1;  
    }  
    exit protocol  
    return;  
}
```

What's a good solution to the mutual exclusion problem?

- A fully satisfactory solution is one that achieves three properties:
 - **Mutual exclusion**: At most one thread is in its critical section at any given time
 - **Freedom from deadlock**: If some threads tries to enter the critical section, some thread will eventually succeed
 - **Freedom from starvation**: Every thread that tries to enter the critical section will eventually succeed
- A good solution should also work for an arbitrary number of threads sharing the same memory
- (Note that freedom from starvation implies freedom from deadlock)

Agenda

- ✓ Concurrent programs
- ✓ Races
- Synchronization problems:
 - ✓ Mutual exclusion
 - Deadlock
 - Starvation and fairness
- Locks

Deadlocks

- A mutual exclusion protocol provides **exclusive access** to shared resources to one thread at a time
 - Threads that try to access the resource when it is not available will have to block and **wait**
 - Mutually dependent waiting conditions may introduce a **deadlock**
- A deadlock is the situation where a group of threads wait forever because each of them is waiting for resources that are held by another thread in the group (circular waiting)

Deadlock: Example

- A **deadlock** is the situation where a group of threads wait forever because each of them is waiting for resources that are held by another thread in the group (circular waiting)
 - A protocol that achieves mutual exclusion but introduces a deadlock
 - **Entry protocol**: wait until all other threads have executed their critical section



Via, resti servita Madama brillante
– E. Tommasi Ferroni, 2012

The dining philosophers


- The dining philosophers is a classic synchronization problem introduced by Dijkstra: It illustrates the problem of deadlocks using a colorful metaphor (by Tony Hoare):
 - Five philosophers are sitting around a dinner table, with a fork in between each pair of adjacent philosophers
 - Each philosopher alternates between thinking (**non-critical section**) and eating (**critical section**)
 - In order **to eat**, a philosopher needs to pick up the **two forks** that lie at the philosopher's left and right sides
 - Since the forks are **shared**, there is a **synchronization** problem between philosophers (threads)



Deadlocking philosophers

- An unsuccessful attempt at solving the dining philosophers problem:

```
entry protocol ( Pk ) {  
    left_fork.acquire();    // pick up left fork  
    right_fork.acquire();   // pick up right fork  
}  
critical section { eat(); }  
exit protocol ( Pk ) {  
    left_fork.release();    // release left fork  
    right_fork.release();   // release right fork  
}
```



- This protocol **deadlocks** if all philosophers get their left forks, and wait forever for their right forks to become available

The Coffman conditions

- Necessary conditions for a deadlock to occur:
 - **Mutual exclusion**: Threads may have exclusive access to the shared resources
 - **Hold and wait**: A thread that may request one resource while holding another resource
 - **No preemption**: Resources cannot forcibly be released from threads that hold them
 - **Circular wait**: Two or more threads form a circular chain where each thread waits for a resource that the next thread in the chain is holding
- Avoiding deadlocks requires to **break one or more** of these conditions

Agenda

- ✓ Concurrent programs
- ✓ Races
- Synchronization problems:
 - ✓ Mutual exclusion
 - ✓ Deadlock
 - Starvation and fairness
- Locks


Starving philosophers

- A solution to the dining philosophers problem that **avoids deadlock** by breaking hold and wait (and thus circular wait): Pick up both forks at once (atomic operation)

```
entry protocol ( Pk ) {
    forks.acquire();    // pick up left and right fork,
                       // atomically
}

critical section { eat(); }

exit protocol ( Pk ) {
    forks.release();    // release left and right fork,
                       // atomically
}
```



This protocol avoids deadlocks, but it may introduce **starvation**: a philosopher may never get a chance to pick up the forks

Starvation

- No deadlocks means that the system makes progress as a whole
- However, some individual thread may still make no progress because it is treated unfairly in terms of access to shared resources
- **Starvation** is the situation where a thread is perpetually denied access to a resource it requests
- **Avoiding starvation** requires some assumption on the scheduler, which has to “give a chance to every thread to execute”

Fairness

■ Weak fairness:

- If a thread continuously requests (that is, requests without interruptions) access to a resource, then access is granted infinitely often

■ Strong fairness:

- If a thread requests access to a resource infinitely often, then access is granted infinitely often


➤ Applied to a scheduler:

- Request = A thread is ready (enabled)
- Fairness = Every thread has a chance to execute

Breaking a circular wait

- Another solution is to pick up first the fork with the lowest id number: This avoids the circular wait because not every philosopher will pick up their left fork first

```
entry protocol ( Pk ) {  
    if (left_fork.id() < right_fork.id()) {  
        left_fork.acquire();  
        right_fork.acquire();  
    } else {  
        right_fork.acquire();  
        left_fork.acquire();  
    }  
}  
critical section { eat(); }  
exit protocol ( Pk ) { /* ... */ }
```



Ordering shared resources and forcing all threads to acquire the resources in order is a common measure to avoid deadlocks

Sequential philosophers

- Another solution that avoids deadlock as well as starvation: a (fair) waiter decides which philosopher eats; the waiter gives permission to eat to one philosopher at a time

```
entry protocol ( Pk ) {  
    while (!waiter.can_eat(k)) {  
        // wait for permission to eat  
    }  
    left_fork.acquire();  
    right_fork.acquire();  
}  
critical section { eat(); }  
exit protocol ( Pk ) { /* ... */ }
```

Having a centralized arbiter avoids deadlocks and starvation, but a waiter who only gives permission to one philosopher a time basically reduces the philosophers to following a sequential order without active concurrency

Agenda

- ✓ Concurrent programs
- ✓ Races
- ✓ Synchronization problems:
 - ✓ Mutual exclusion
 - ✓ Deadlock
 - ✓ Starvation and fairness
- Locks

Lock objects

- A lock is a data structure with interface:

```
interface Lock {  
    void lock();        // acquire lock  
    void unlock();      // release lock  
}
```

- Several threads share the same object lock of type Lock
 - Multiple threads calling lock.lock() results in exactly one thread T acquiring the lock:
 - T's call lock.lock() returns: T is holding the lock
 - Other threads block on the call lock.lock(), waiting for the lock to
 - Become available
 - A thread T that is holding the lock calls lock.unlock() to release the lock:
 - T's call lock.unlock() returns; the lock becomes available
 - Another thread waiting for the lock may succeed in acquiring it
- Locks are also called **mutexes**; they guarantee mutual exclusion

Using locks

- With lock objects the entry/exit protocols are trivial:
 - **Entry protocol**: call `lock.lock()`
 - **Exit protocol**: call `lock.unlock()`

<code>int counter = 0; Lock lock = new Lock();</code>	
<code>thread t</code>	<code>thread u</code>
<code>int cnt;</code>	<code>int cnt;</code>
<code>lock.lock();</code>	<code>lock.lock();</code>
<code> cnt = counter;</code>	<code> cnt = counter;</code>
<code> counter = cnt + 1;</code>	<code> counter = cnt + 1;</code>
<code>lock.unlock();</code>	<code>lock.unlock();</code>

- The implementation of the Lock interface should guarantee mutual exclusion, deadlock freedom, and starvation freedom

Using locks in Java

```
// package with lock-related classes
import java.util.concurrent.locks.*;

// shared with other synchronizing threads
Lock lock;

while(true) {
    lock.lock(); // entry protocol
    try {
        // critical section
        // mutual exclusion is guaranteed
        // by the lock protocol
    } finally {    // lock released even if
                   // an exception is thrown in the critical section
        lock.unlock(); // exit protocol
    }
}
```

Counter with mutual exclusion

```
public class LockedCounter extends CCounter
{
    @Override
    public void run() {
        lock.lock();
        try {
            // int cnt = counter;
            // counter = counter + 1;
            super.run();
        } finally {
            lock.unlock();
        }
    }

    // shared by all threads working on this object
    private Lock lock = new ReentrantLock();
}
```

Built-in locks in Java

- Every object in Java has an implicit lock, which can be accessed using the keyword `synchronized`:

- Whole method locking
(`synchronized` methods):

```
synchronized T m() {  
    // the critical section  
    // is the whole method  
    // body  
}
```

- Every call to `m` implicitly:
 - Acquires the lock
 - Executes `m`
 - Releases the lock

- Block locking
(`synchronized` block):

```
synchronized (this) {  
    // the critical section  
    // is the block's content  
}
```

- Every execution of the block implicitly:
 - Acquires the lock
 - Executes `m`
 - Releases the lock

Counter with mutual exclusion: with synchronized

■ SyncCounter class:

```
public class SyncCounter
    extends CCounter
{
    @Override
    public synchronized
    void run() {
        // int cnt = counter;
        // counter = counter + 1;
        super.run();
    }
}
```

■ SyncBlockCounter class:

```
public class SyncBlockCounter
    extends CCounter
{
    @Override
    public void run() {
        synchronized(this) {
            // int cnt = counter;
            // counter = counter + 1;
            super.run();
        }
    }
}
```


Lock implementations in Java

- The most common implementation of the Lock interface in Java is class ReentrantLock
 - **Mutual exclusion:**
 - ReentrantLock guarantees mutual exclusion
 - **Starvation:**
 - ReentrantLock does not guarantee freedom from starvation by default
 - However, calling the constructor with new ReentrantLock(true) “favors granting access to the longest-waiting thread”
 - This still does not guarantee that thread scheduling is fair
 - **Deadlocks:**
 - One thread will succeed in acquiring the lock
 - However, deadlocks may occur in systems that use multiple locks (remember the dining philosophers)

Built-in lock implementations in Java

- The built-in locks — used by synchronized methods and blocks — have the same behavior as the explicit locks of `java.util.concurrent.locks`
- Built-in locks, as well as all lock implementations in `java.util.concurrent.locks`, are re-entrant: A thread holding a lock can lock it again without causing a deadlock

