

Semaphores

Dr. José Luis Zechinelli Martini

joseluis.zechinelli@udlap.mx

LDS-1101

Based on the course of Professor Carlo A. Furia, Chalmers University of Technology – University of Gothenburg

Agenda

- Semaphores:
 - Semaphores for permissions
 - Mutual exclusion for two processes
 - Weak vs. strong semaphores
 - Invariants
 - Binary semaphores
 - Using semaphores in Java

Semaphores

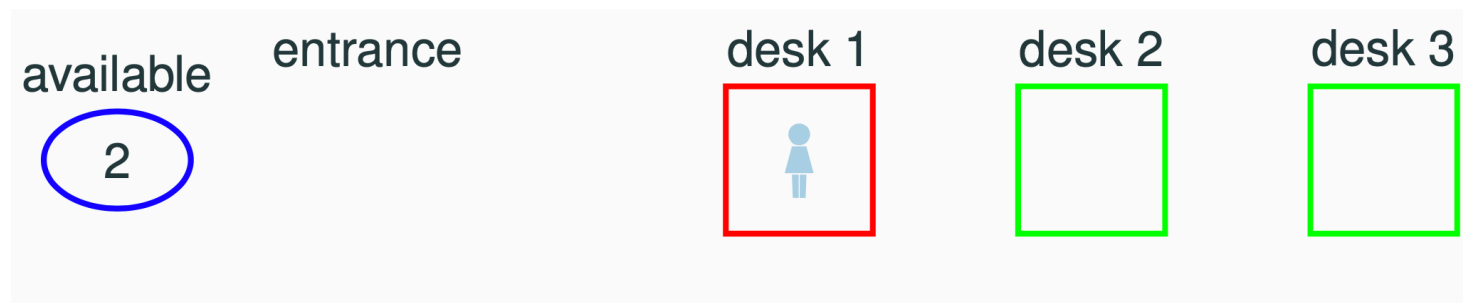
- A (general/counting) semaphore is a data structure with interface:

```
interface Semaphore {  
    int count();    // current value of counter  
    void up();      // increment counter  
    void down();    // decrement counter  
}
```

- Several threads share the same object **sem** of type Semaphore:
 - Initially count is set to a nonnegative value C , *i.e.*, the capacity
 - A call to `sem.up()` atomically increments count by one
 - A call to `sem.down()` waits until count is positive, and then atomically decrements count by one

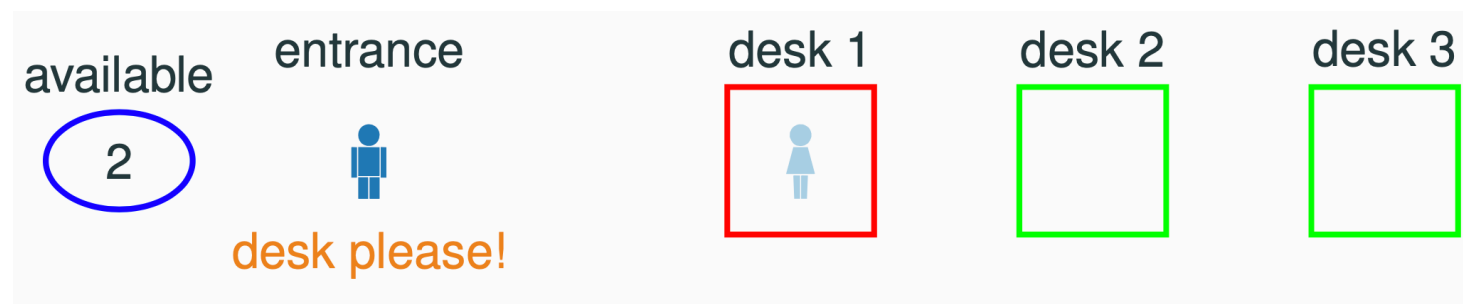
Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**




Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



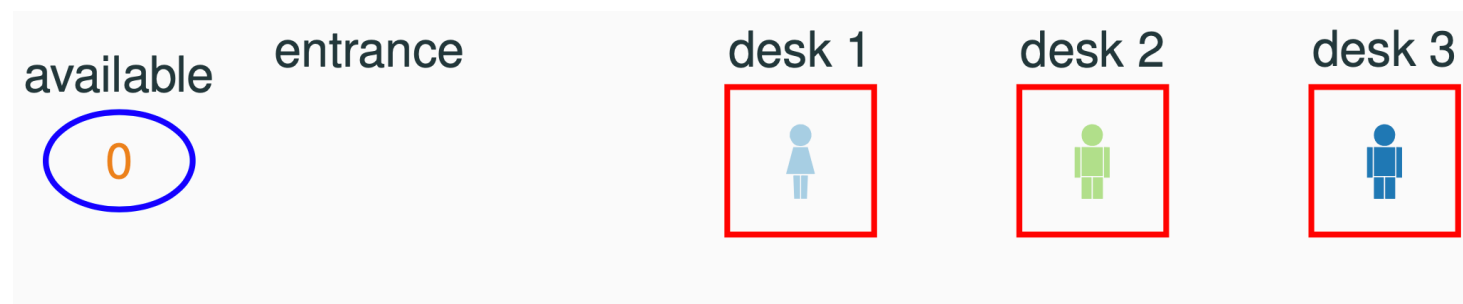
Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



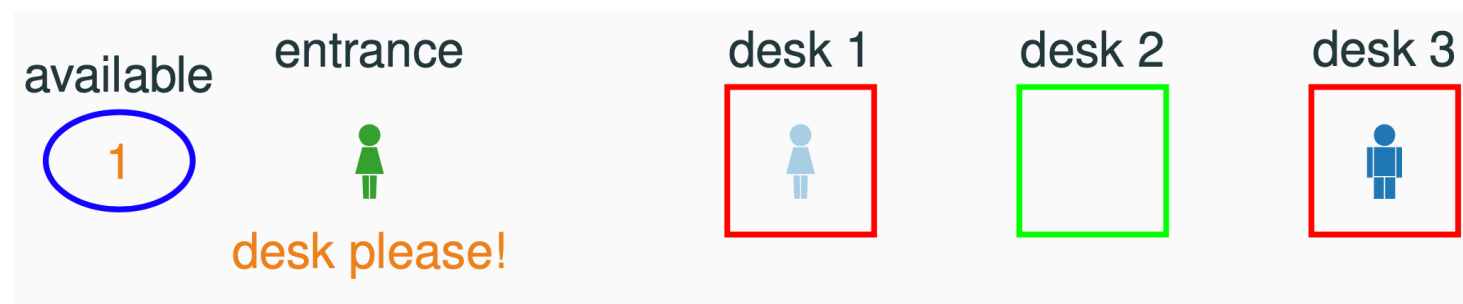
Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



Semaphores for permissions

- A semaphore is often used to regulate access permits to a finite number of resources:
 - The capacity C is the number of **initially available** resources
 - Up (also called signal) **releases** a resource, which becomes available
 - Down (also called wait) **acquires** a resource if it is available
- Example: **Hot Desks**



Agenda

- Semaphores:
 - ✓ Semaphores for permissions
 - Mutual exclusion for two processes
 - Weak vs. strong semaphores
 - Invariants
 - Binary semaphores
 - Using semaphores in Java

Mutual exclusion for two processes with semaphores

- With semaphores the entry/exit protocols are trivial:
 - Initialize semaphore to 1
 - **Entry protocol**: call `sem.down()`
 - **Exit protocol**: call `sem.up()`



- The implementation of the Semaphore interface guarantees mutual exclusion, deadlock freedom, and starvation freedom

Weak vs. strong semaphores

- Every implementation of semaphores should guarantee the atomicity of the up and down operations, as well as deadlock freedom:
 - For threads only sharing one semaphore, deadlocks may still occur if there are other synchronization constraints
- Fairness is optional:
 - **Weak semaphore**: Threads waiting to perform down are scheduled non-deterministically
 - **Strong semaphore**: Threads waiting to perform down are scheduled fairly in FIFO (First In First Out) order

Agenda

- Semaphores:
 - ✓ Semaphores for permissions
 - ✓ Mutual exclusion for two processes
 - ✓ Weak vs. strong semaphores
 - Invariants
 - Binary semaphores
 - Using semaphores in Java

Invariants

- An object's invariant is a property that always holds between calls to the object's methods:
 - The invariant holds initially (when the object is created)
 - Every method call starts in a state that satisfies the invariant
 - Every method call ends in a state that satisfies the invariant
- For example: A bank account that cannot be overdrawn has an invariant ($\text{balance} \geq 0$)

```
class BankAccount {  
    int balance = 0;  
    void deposit(int amount)  
        { if (amount > 0) balance += amount; }  
    void withdraw(int amount)  
        { if (amount > 0 && balance > amount) balance -= amount; }  
}
```

Invariants in pseudo-code

- We occasionally annotate classes with invariants using the pseudo-code keyword `invariant`
- Note that `invariant` is not a valid Java keyword — that is why we highlight it in a different color — but we will use it whenever it helps make more explicit the behavior of classes

```
class BankAccount {  
    int balance = 0;  
    void deposit(int amount)  
        { if (amount > 0) balance += amount; }  
    void withdraw(int amount)  
        { if (amount > 0 && balance > amount) balance -= amount; }  
    invariant { balance >= 0; } // not valid Java code  
}
```

Invariants of semaphores

- A semaphore object with capacity C satisfies the invariant:

```
interface Semaphore {  
    int count();  
    void up();  
    void down();  
    invariant {  
        count() >= 0;  
        count() == C + #up - #down;  
    }  
}
```

Number of calls to down

Number of calls to up

- Invariants characterize the behavior of an object, and are very useful for proofs

Agenda

- Semaphores:
 - ✓ Semaphores for permissions
 - ✓ Mutual exclusion for two processes
 - ✓ Weak vs. strong semaphores
 - ✓ Invariants
 - Binary semaphores
 - Using semaphores in Java

Binary semaphores

- A semaphore with capacity 1 and operated such that `count()` is always at most 1 is called a binary semaphore:

```
interface BinarySemaphore extends Semaphore {  
    invariant  
    { 0<=count()<=1;  
      count() == C + #up - #down; }  
}
```

- Mutual exclusion uses a binary semaphore:

```
Semaphore sem = new Semaphore(1); // shared by all threads  
  
thread t  
  
sem.down();  
    // critical section  
sem.up();
```

- If the semaphore is strong this guarantees starvation freedom

Binary semaphores vs. locks

- Binary semaphore are very similar to locks with one difference:
 - In a lock, only the thread that decrements the counter to 0 can increment it back to 1
 - In a semaphore, a thread may decrement the counter to 0 and then let another thread increment it to 1
- Thus (binary) semaphores support transferring of permissions

Agenda

- Semaphores:
 - ✓ Semaphores for permissions
 - ✓ Mutual exclusion for two processes
 - ✓ Weak vs. strong semaphores
 - ✓ Invariants
 - ✓ Binary semaphores
 - Using semaphores in Java

Using semaphores in Java

- Method acquire may throw an InterruptedException (catch or propagate):

```
package java.util.concurrent;

public class Semaphore {
    Semaphore(int permits);           // initialize with capacity permits
    Semaphore(int permits, boolean fair); // fair  $\Leftrightarrow$  fair semaphore
        // fair == true also called 'strong' semaphore
        // fair == false also called 'weak' semaphore
    void acquire();                   // corresponds to down
    void release();                   // corresponds to up
    int availablePermits();           // corresponds to count
}
```